



AI VISIBILITY MEASUREMENT - FOUR ENGINES

Credible PR

Prepared for Max Muir - collected 2026-09-08, one night, one prompt bank

0 of 350

Credible PR appears in 0 of 350 AI responses, across four engines.
Claude 0/88 - ChatGPT 0/90 - Gemini 0/88 - Perplexity 0/84

BANK pr_agency_v1 - 30 PROMPTS, 5 FAMILIES - 3 RUNS/PROMPT - RUNG R1
CLAUDE MODEL claude-sonnet-5 - TEMP omitted(rejected-or-na) - CLEAN 88 of 90 - SHA-256 5ac8d9d8...f9bed
CHATGPT MODEL chat-latest - TEMP 0.0 - CLEAN 90 of 90 - SHA-256 82660a1f...be696
GEMINI MODEL gemini-3.8-flash - TEMP 0.0 - CLEAN 88 of 90 - SHA-256 81fc2f91...0707f
PERPLEXITY MODEL sonar - TEMP 0.0 - CLEAN 84 of 90 - SHA-256 442259f4...70615

ChatGPT measured via the OpenAI consumer-chat alias 'chat-latest'; the concrete snapshot is not disclosed by the API at call time. Claude rejects the temperature parameter; it is omitted and declared. Gemini free tier: inputs and outputs may be used by Google to improve its models. Copilot excluded (no public API). Grok excluded (<3% share, REPORTED 2026-09-08). All engines are developer-API surfaces; consumer apps add search and memory.

Executive brief

WHAT WAS MEASURED. 30 prompts a prospective client might put to an AI assistant about PR firms, each run 3 times against four engines - Claude, ChatGPT, Gemini, Perplexity - on 2026-09-08. 360 responses attempted, 350 clean; 10 lost to transport errors and excluded from every rate.

WHAT WAS FOUND. Credible PR is named in none of the 350 responses, on any engine, in any prompt family, including every buyer-intent prompt. The same measurement shows which firms are named (page 3), so the zero is a property of the answers, not of the instrument. The absence holds on four independent surfaces measured the same night.

WHAT ELSE WAS FOUND. The engines do not agree on who leads. Claude's most-named firms differ from ChatGPT's, and Gemini concentrates heavily on one incumbent. Which AI a buyer uses changes which firms they hear about - a per-engine question no single-engine measurement can answer.

WHAT THIS DOES NOT SAY. This measures what models answer, not what people ask; it carries no information about search volume or revenue. Name extraction is pattern-matching against a declared list of 23 firms; its error rate is unmeasured until our agreement study publishes. Per-engine samples are small (84-90) and cross-engine gaps mostly do not separate statistically - the intervals on page 3 say so.

Three actions

1. Baseline each client market on these four engines.

Is the client winning where its buyers actually ask? ~\$3 per market.


2. Re-measure 30 days after the next placement lands.

Does a placement move any engine's answers? The pilot's first join.

3. Prioritise the engine your buyers use.

Gemini's answers differ most; ChatGPT carries the most share.

Who each engine names - and who it doesn't



Overall presence per engine: each engine's three most-named firms, then Credible PR in red. Whiskers are 95% Wilson intervals. Denominators are clean responses per engine. Where intervals overlap, no separation is claimed - at these sample sizes most cross-engine gaps do not separate.

All candidates, all engines

candidate (overall presence)	Claude n=88	ChatGPT n=90	Gemini n=88	Perplexity n=84
Edelman	13	12	24	10
5WPR	0	19	16	13
Otter PR	0	19	18	9
Bospar	17	9	13	3
Weber Shandwick	13	3	11	9
FINN Partners	3	15	0	8
LaunchSquad	1	2	14	3
Ketchum	9	3	0	5
FleishmanHillard	8	1	1	7
SourceCode Communications	0	14	1	0
Red Banyan	0	12	1	0
The Lede Company	3	6	2	0
Highwire	1	4	4	1
Method Communications	5	1	1	1
SimplyBe. Agency	0	3	0	4
Burson	0	1	0	5
Clarity PR	0	0	2	0
Channel V Media	0	0	2	0
Reputation Ink	0	1	0	0
Bollare	0	1	0	0
Reputation Rhino	0	0	0	0
Credible PR	0	0	0	0
Brandstyle	0	0	0	0

Buyer-intent prompts specifically - the questions asked at the moment of hiring: Credible PR: Claude 0/16 - ChatGPT 0/18 - Gemini 0/18 - Perplexity 0/17. Per-family tables with intervals for every engine are reproducible from the evidence files.

Transport errors were not evenly distributed (Perplexity's fell in one family), so every rate in the evidence prints against its own family denominator, never a pooled one.

What this report refuses to claim

NO COMPOSITE SCORE.

A single blended number hides which component moved. Every rate stands alone with its denominator and interval.

NO POOLING ACROSS ENGINES.

Four engines are four surfaces. Rates are never averaged across them; the per-engine number is the measurement.

NO ORDERING CLAIM.

No order-bias test has been run on this bank, so no ordering claim is made.

NO EXTRACTION-ERROR FIGURE.

Name extraction is pattern-matching against the declared list of 23. Its agreement with human judgement is being measured, with the labelled set to be published; until then no figure is quoted.

NO CROSS-ENGINE SEPARATION AT THESE SAMPLES.

At 84-90 responses per engine, most differences between engines sit inside overlapping intervals. Divergence is reported as observed, not as established.

NO DEMAND OR CAUSAL CLAIM.

This measures what models answer, not what people ask, and says nothing about what moves these rates - that is the pilot's question.

Each refusal names what would unlock the claim. Ask.

Check this yourself

Four evidence files ship with this report - every response, its request payload, its hash - plus `verify_evidence.py`, a checker that needs only standard Python. No install, nothing from us. One command per file:

```
python verify_evidence.py multi_anthropic_run3_20260908.jsonl \
--expect 5ac8d9d86829da414c3a2baaf122bb504cda7c551a562161846e72ab92cf9bed

python verify_evidence.py multi_openai_run3_20260908.jsonl \
--expect 82660a1f5d5e9782ac1fe0e03e6642214af2ae98d136badc8da1bd1a012be696

python verify_evidence.py multi_gemini_run3_20260908.jsonl \
--expect 81fc2f91615cfe583592a4871129a2f8be05596854af5b15c9d9fbf69190707f

python verify_evidence.py multi_perplexity_run3_20260908.jsonl \
--expect 442259f4cf0fccdde35093602296ef3a556d0ee0943bf4f13bb57ecd263870615
```

Expected last line each time: VERDICT PASS.

WHAT PASS MEANS: the stored text and the stored hashes are consistent with each other. The file has not been edited since it was written.

WHAT PASS DOES NOT MEAN: that the responses came from the engine named in the payload, or that the run happened as described. A hash cannot establish either. Ask for the console receipt and the probe record.

If you run this and any verdict is anything but PASS, that is a finding - send it back.