



AI VISIBILITY MEASUREMENT - FOUR ENGINES

Credible PR

Prepared for Max Muir - collected 2026-09-08, one night, one prompt bank

0 of 350

Credible PR appears in 0 of 350 clean AI responses (360 attempted), across four engines.

Claude 0/88 - ChatGPT 0/90 - Gemini 0/88 - Perplexity 0/84

0 of 120 prompts on any engine - the 3 runs per prompt are near-replicates, so the prompt is the sampling unit
a zero can rule out at this sample (95% upper bound on the true prompt rate): Claude <= 11.4% - ChatGPT <= 11.4% - Gemini <= 11.4% - Perplexity <= 11.4%

*Why this number matters: when a prospect asks an AI to recommend PR firms, the answer draws from a set of names.
This page states whether Credible PR is in that set today - the baseline any campaign would be measured against.*

BANK pr_agency_v1 - 30 PROMPTS, 5 FAMILIES - 3 RUNS/PROMPT - RUNG R1
CLAUDE MODEL claude-sonnet-5 - TEMP omitted(rejected-or-na) - CLEAN 88 of 90 - SHA-256 5ac8d9d8...f9bed
CHATGPT MODEL chat-latest - TEMP 0.0 - CLEAN 90 of 90 - SHA-256 82660a1f...be696
GEMINI MODEL gemini-3.8-flash - TEMP 0.0 - CLEAN 88 of 90 - SHA-256 81fc2f91...0707f
PERPLEXITY MODEL sonar - TEMP 0.0 - CLEAN 84 of 90 - SHA-256 442259f4...70615

ChatGPT measured via the OpenAI consumer-chat alias 'chat-latest'; the concrete snapshot is not disclosed by the API at call time. Claude rejects the temperature parameter; it is omitted and declared. Gemini free tier: inputs and outputs may be used by Google to improve its models; this bank contains only generic prospect questions, no client data. Copilot excluded (no public API). Grok excluded (<3% share, REPORTED 2026-09-08). All engines are developer-API surfaces: consumer apps add search and memory.

Executive brief

WHY THIS PAGE MATTERS - any tool can print a zero. A zero is only meaningful if the reader knows what was asked, how many times, on which engines, and what was excluded - so this page states the method before the finding.

WHAT WAS MEASURED. 30 prompts a prospective client might put to an AI assistant about PR firms, each run 3 times against four engines - Claude, ChatGPT, Gemini, Perplexity - on 2026-09-08. 360 responses attempted, 350 clean; 10 lost to transport errors and excluded from every rate.

WHAT WAS FOUND. Credible PR is named in none of the 350 responses, on any engine, in any prompt family, including every buyer-intent prompt. Comparably small firms do register (page 4). The matcher is deliberately biased toward finding the subject: it matches case-insensitively, so even the ordinary phrase 'credible PR' in a sentence counts as a hit. It still found nothing, on four independent surfaces measured the same night.

WHAT ELSE WAS FOUND. The engines do not agree on who leads. Claude's most-named firms differ from ChatGPT's, and Gemini concentrates heavily on one incumbent. Which AI a buyer uses changes which firms they hear about - a per-engine question no single-engine measurement can answer. Pages 5 and 6 break this down by prompt family and by position in the answer.

WHAT THIS DOES NOT SAY. This measures what models answer, not what people ask; it carries no information about search volume or revenue. Name extraction is pattern-matching against a declared list of 23 firms; its error rate is unmeasured until the agreement study publishes. Per-engine samples are small (84-90) and cross-engine gaps mostly do not separate statistically - the intervals on page 3 say so.

Three actions

1. Baseline each client market on these four engines.

Is the client winning where its buyers actually ask? ~\$3 per market [EST: unit cost measured on Claude, extrapolated to the other engines at current API prices] - negligible next to one new client.

2. Re-measure 30 days after the next placement lands.


Does a placement move any engine's answers? The pilot's first join.

3. Prioritise the engine your buyers use.

Gemini's answers differ most; ChatGPT carries the most share [REPORTED 2026-09-08].

Who each engine names - and who it doesn't

WHY THIS PAGE MATTERS - the four engines disagree about who leads, and every buyer uses one engine at a time - so an average across them would describe a market nobody is in. Rates are counted per PROMPT, not per response, because the 3 runs of one prompt are near-replicates: counting them separately would triple the sample without adding information. The whiskers show where the true rate plausibly sits at this sample size.



Share of PROMPTS naming each firm - a prompt counts if any of its 3 runs names it. Each engine's three most-named firms, then Credible PR in red. Whiskers are 95% Wilson intervals at the prompt level, the honest sampling unit given near-replicate runs. Where intervals overlap, no separation is claimed.

All candidates, all engines

WHY THIS PAGE MATTERS - measuring the subject alone cannot distinguish 'these models never name boutiques' from 'these models do not name this boutique'. The 22 competitor rows are the control: small firms that DO register prove the instrument finds small firms.

candidate (overall presence)	Claude n=88	ChatGPT n=90	Gemini n=88	Perplexity n=84
Edelman	13	12	24	10
5WPR	0	19	16	13
Otter PR	0	19	18	9
Bospar	17	9	13	3
Weber Shandwick	13	3	11	9
FINN Partners	3	15	0	8
LaunchSquad	1	2	14	3
Ketchum	9	3	0	5
FleishmanHillard	8	1	1	7
SourceCode Communications	0	14	1	0
Red Banyan	0	12	1	0
The Lede Company	3	6	2	0
Highwire	1	4	4	1
Method Communications	5	1	1	1
SimplyBe. Agency	0	3	0	4
Burson	0	1	0	5
Clarity PR	0	0	2	0
Channel V Media	0	0	2	0
Reputation Ink	0	1	0	0
Bollare	0	1	0	0
Reputation Rhino	0	0	0	0
Credible PR	0	0	0	0
Brandstyle	0	0	0	0

3 of 23 declared candidates scored zero on every engine: Reputation Rhino, Credible PR, Brandstyle. The subject's zero is one of 3, found by a list declared before the run - not an artifact of a list built to find it.

Buyer-intent prompts specifically - the questions asked at the moment of hiring - with the 95% upper bound each zero supports:
Credible PR: Claude 0/16 ($\leq 19\%$) - ChatGPT 0/18 ($\leq 18\%$) - Gemini 0/18 ($\leq 18\%$) - Perplexity 0/17 ($\leq 18\%$). Small samples give wide bounds; printing them wide is the point.

Transport errors were not evenly distributed (Perplexity's fell in one family), so every rate in the evidence prints against its own family denominator, never a pooled one.

How much the engines agree

Overlap of each pair's five most-named firms: Clau-Chat 1/5 - Clau-Gemi 3/5 - Clau-Perp 3/5 - Chat-Gemi 3/5 - Chat-Perp 3/5 - Gemi-Perp 3/5. Fewer shared names means a buyer's engine choice changes which firms they hear about.

Three of the buyer-intent prompts, verbatim

"I want to hire a PR agency to get me into national media. Who should I contact?"

"Recommend a PR firm for a solo professional with a three thousand dollar monthly budget."

"Which PR agency should I hire to get featured in USA Today?"

The category structure of the bank is public; the full instantiated prompt set ships to the audited party only, so it cannot be optimised against.

Five prompt families, four engines

WHY THIS PAGE MATTERS - a zero on 'best PR firms' and a zero on 'should I hire a firm or do it myself' are different business problems. The five families are declared prompt text, not simulated audiences - and transport errors clustered by family (a measured property of this run), so every cell prints its own denominator.

family (prompt-level)	Claude	ChatGPT	Gemini	Perplexity
buyer_intent				
Credible PR	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)
most-named	Bospar 1/6	Otter PR 4/6	Otter PR 4/6	5WPR 3/6
category				
Credible PR	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)
most-named	Bospar 6/6	5WPR 5/6	Edelman 5/6	Otter PR 3/6
comparison				
Credible PR	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)
most-named	Edelman 2/6	5WPR 2/6	Edelman 2/6	Edelman 2/6
long_tail				
Credible PR	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)
most-named	5WPR 0/6	5WPR 0/6	Edelman 1/6	5WPR 0/6
problem				
Credible PR	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)	0/6 (<=39%)
most-named	5WPR 0/6	5WPR 0/6	5WPR 0/6	5WPR 0/6

Reading the grid: the subject's cell shows presence with the 95% Wilson upper bound its zero (or count) supports at that family's prompt count; the second line names the family's most-named firm on that engine. Family-level samples are small (5-6 prompts) - the bounds are wide, and printing them wide is the point. Per-family tables for all 23 candidates are reproducible from the evidence files.

Named early, or merely named

WHY THIS PAGE MATTERS - being one name in a list of ten is not the same as being the first name a buyer reads. This page counts, per prompt, whether a firm appears among the FIRST THREE distinct candidate names in the answer text, by first occurrence. The candidate list holds 23 names, above the five-name floor this metric requires - a shorter list would make first-three collapse into mention.

Claude (30 prompts, claude-sonnet-5)

Edelman	7/30	95% CI	12- 41%
Weber Shandwick	7/30	95% CI	12- 41%
Bospar	5/30	95% CI	7- 34%
Credible PR	0/30	95% CI	0- 11%

ChatGPT (30 prompts, chat-latest)

5WPR	9/30	95% CI	17- 48%
Otter PR	6/30	95% CI	10- 37%
SourceCode Communications	6/30	95% CI	10- 37%
Credible PR	0/30	95% CI	0- 11%

Gemini (30 prompts, gemini-3.8-flash)

Edelman	10/30	95% CI	19- 51%
Otter PR	8/30	95% CI	14- 44%
5WPR	7/30	95% CI	12- 41%
Credible PR	0/30	95% CI	0- 11%

Perplexity (30 prompts, sonar)

5WPR	6/30	95% CI	10- 37%
Edelman	4/30	95% CI	5- 30%
Otter PR	3/30	95% CI	3- 26%
Credible PR	0/30	95% CI	0- 11%

A prompt counts for a firm if ANY of its 3 runs places that firm in the answer's first three distinct candidate names. This is an observation about where names sit in these stored answers - it is not a validated ordering metric; see the refusal on page 8 for exactly what is and is not claimed.

What the engines actually said

WHY THIS PAGE MATTERS - every number in this report is computed from stored response text, and this page shows some of that text. Selection is by fixed rule, not curation: for each engine, its most-named firm, then the first buyer-intent row in file order naming that firm (first clean row anywhere if buyer-intent has none). Verbatim model text sits between the EXHIBIT markers; it is the engines' language, not aivis's. Whitespace is normalised and non-Latin characters replaced for layout - the byte-exact text lives in the evidence file at the row hash printed above each exhibit.

Claude - most-named firm Bospar - first buyer-intent row naming it

prompt PR-INT05 - run 1 - row sha256 8fcc39b48050a144ba0fef43c4e17db5880667c842b96887a4e511dfb475b14f

--- EXHIBIT BEGIN ---

Building Authority as an Independent Consultant: PR Strategy Guide I can't recommend a single "best" agency since that depends heavily on your budget, industry, and goals?but I can give you a practical framework for finding the right fit. ## Questions to Answer First - **Budget reality**: Boutique PR firms for solo consultants typically run \$3-10K/month; top-tier agencies want \$15K+/month - **Niche specificity**: Do you n [...truncated for layout - full text in the evidence file at the row hash above]

--- EXHIBIT END ---

ChatGPT - most-named firm 5WPR - first buyer-intent row naming it

prompt PR-INT01 - run 0 - row sha256 40abffa9ab12b828b8f529f29964ac36b78d858922c69e9091c7fa2b83817895

--- EXHIBIT BEGIN ---

For national U.S. media?think The New York Times, WSJ, Forbes, Bloomberg, CNBC, CNN, Fast Company, major podcasts, and broadcast?the right agency depends heavily on what you want publicity for. Some strong firms to consider: - DKC ? large national agency with deep media relationships; good for consumer brands, executives, entertainment, tech, and corporate stories. dkcnnews.com - Sunshine Sachs Morgan & Lylis ? high-profile pe [...truncated for layout - full text in the evidence file at the row hash above]

--- EXHIBIT END ---

Gemini - most-named firm Edelman - first buyer-intent row naming it

prompt PR-INT01 - run 0 - row sha256 a5568b5217ac7853d332030c92d8711324b02142b00272c250d0103a9c31e58b

--- EXHIBIT BEGIN ---

To give you the best recommendation, it helps to know your industry, but if your goal is strictly **top-tier national media** (e.g., *The New York Times, The Wall Street Journal, TODAY, Good Morning America, Forbes, CNBC, Bloomberg*), you need an agency with strong editorial relationships and a proven track record. Here is a breakdown of the top PR agencies known for landing national coverage, categorized by industry focus: [...truncated for layout - full text in the evidence file at the row hash above]

--- EXHIBIT END ---

Perplexity - most-named firm 5WPR - first buyer-intent row naming it

prompt PR-INT04 - run 0 - row sha256 70928a751809caedfa7b7bdbd042a2f1bb9de83747e71f3898aaa7fe2043b929

--- EXHIBIT BEGIN ---

If you want a **PR agency specifically for an aesthetic clinic**, the strongest matches from the results are **Cosmetic PR**, **Kendrick PR**, **Illumination PR**, **Society22 PR**, and **5WPR**.[20][3][6][4][18] My top recommendation depends on your clinic type: - **Best boutique choice:** **Cosmetic PR** ? a full-service boutique agency that specializes in **beauty, medical aesthetics, and wellness** brands.[20] - **Best f [...truncated for layout - full text in the evidence file at the row hash above]

--- EXHIBIT END ---

What this report refuses to claim

WHY THIS PAGE MATTERS - every claim a report declines to make is a claim a reader might otherwise assume. Naming the limits is what makes the numbers inside them usable - and each refusal below names what would unlock the claim.

NO COMPOSITE SCORE.

A single blended number hides which component moved. Every rate stands alone with its denominator and interval.

NO POOLING ACROSS ENGINES.

Four engines are four surfaces. Rates are never averaged across them; the per-engine number is the measurement.

POSITION-IN-ANSWER: OBSERVED, NOT VALIDATED.

Page 6 counts who appears in an answer's first three names. That is an observation about these stored responses, not a validated ordering metric: no order-flip test has been run on this bank, so no claim is made that the ordering is stable or meaningful. The order-flip delta is the unlock.

NO EXTRACTION-ERROR FIGURE.

Name extraction is pattern-matching against the declared list of 23. Its agreement with human judgement is being measured, with the labelled set to be published; until then no figure is quoted.

NO SOURCE OR CITATION TABLE.

Which domains the engines cite is the most actionable data in this category, and this run cannot show it honestly: the v1 runner stored response text only, and that text contains 16 URLs across all 350 clean responses [computed at emit time]. Unlocked by capturing each engine's citation fields from the next run onward.

NO SENTIMENT SCORE.

Sentiment classifiers ship unvalidated across this category. None exists here with a published agreement figure, so no sentiment number ships. The unlock is the same agreement study that governs extraction.

NO CROSS-ENGINE SEPARATION AT THESE SAMPLES.

At 84-90 responses per engine, most differences between engines sit inside overlapping intervals. Divergence is reported as observed, not as established.

NO DEMAND OR CAUSAL CLAIM.

This measures what models answer, not what people ask, and says nothing about what moves these rates - that is the pilot's question.

NO CROSS-VERSION COMPARISON.

These baselines are pinned to the model versions on page 1. Providers retire models within quarters; every re-measure declares its own surface, and rates are never compared across model versions as if they were one instrument.

Each refusal names what would unlock the claim. Ask.

Check this yourself

WHY THIS PAGE MATTERS - reports in this category ask to be trusted. This one ships its raw material instead. If the check fails, the report is wrong - that is the point of the check.

Four evidence files ship with this report - every response, its request payload, its hash - plus `verify_evidence.py`, a checker that needs only standard Python. No install, nothing from us. One command per file:

```
python verify_evidence.py multi_anthropic_run3_20260908.jsonl \  
--expect 5ac8d9d86829da414c3a2baaf122bb504cda7c551a562161846e72ab92cf9bed  
python verify_evidence.py multi_openai_run3_20260908.jsonl \  
--expect 82660a1f5d5e9782ac1fe0e03e6642214af2ae98d136badc8da1bd1a012be696  
python verify_evidence.py multi_gemini_run3_20260908.jsonl \  
--expect 81fc2f91615cfe583592a4871129a2f8be05596854af5b15c9d9fbf69190707f  
python verify_evidence.py multi_perplexity_run3_20260908.jsonl \  
--expect 442259f4cf0fcdde35093602296ef3a556d0ee0943bf4f13bb57ecd263870615
```

Expected last line each time: VERDICT PASS.

WHAT PASS MEANS: the stored text and the stored hashes are consistent with each other. The file has not been edited since it was written.

WHAT PASS DOES NOT MEAN: that the responses came from the engine named in the payload, or that the run happened as described. A hash cannot establish either. Ask for the console receipt and the probe record.

If you run this and any verdict is anything but PASS, that is a finding - send it back.