Which restaurants AI assistants name: New York City

10 neighbourhoods, 4 assistants, 400 answers. Collected 2026-10-07 02:26 to 2026-10-07 03:52 UTC.

PRELIMINARY - NAMES NOT CHECKED. Every name in this report was pulled out of the answers by software. None has been checked against the real world, some entries are not restaurants, and some restaurants appear under two spellings. The decision rule has not been evaluated. Not for sending to any restaurant.

```
ASSISTANTS 4 | PROMPTS 20 | RUNS PER PROMPT 5 | CLEAN ANSWERS 400 of 400 | ERROR ROWS 0
MODELS Claude: claude-sonnet-5 | Gemini: gemini-3.8-flash | ChatGPT: gpt-4o-2024-08-06 | Perplexity: perplexity/sonar
SEARCH TOOL OFFERED on 400 of 400 answers | ANSWERS CITING A SOURCE 304 of 400 | OUTPUT CAP 8192 tokens
PROMPT BANK nyc_restaurant_prompts_v1 | SURFACE developer API, not the consumer apps
RUNG R1 (ran once on the operator's machine) | NAME EXTRACTION UNVALIDATED_PARSE_PENDING_KAPPA
RUN ID 3dadeb35b62fe85f
```

What was asked

Twenty everyday questions about where to eat (best overall, date night, brunch, a group dinner, cheap eats and so on), two for each of 10 neighbourhoods. No question names a restaurant. Each question was put to each assistant five times.

What came back

- 6 names appeared in at least three quarters of the answers for its own neighbourhood: Le Bernardin (40 of 40, Midtown Manhattan); Taverna Kyclades (40 of 40, Astoria); Buvette (36 of 40, West Village); Beso (35 of 40, St. George); Zero Otto Nove (35 of 40, Arthur Avenue); Enoteca Maria (31 of 40, St. George).
- 124 names were given by at least 3 of the 4 assistants for the same neighbourhood - before checking, so the figure includes second spellings and some entries that are not restaurants.
- The ten most-named account for 10.5% (324 of 3088) of all mentions: no single restaurant dominates across neighbourhoods.
- Answers that cite at least one source: Claude 100 of 100, Gemini 35 of 100, ChatGPT 69 of 100, Perplexity 100 of 100. An answer with no citation shows no sign of a live search.

Most frequently named, by neighbourhood

Each neighbourhood has 40 answers. The interval is a 95% Wilson interval on the share of those answers. This is not an ordering: two names whose intervals overlap are not placed one above the other, here or anywhere.

Neighbourhood	Name as written by the assistants	Answers	95% interval	Assistants	Named first
Astoria	Taverna Kyclades	40 of 40	91-100%	4 of 4	27
	Bahari Estiatorio	17 of 40	29-58%	4 of 4	1
	Milkflower	17 of 40	29-58%	3 of 4	0
	Ruta Oaxaca	14 of 40	22-50%	3 of 4	0
Midtown Manhattan	Le Bernardin	40 of 40	91-100%	4 of 4	25
	Keens Steakhouse	27 of 40	52-80%	4 of 4	0
	Midtown East *	24 of 40	45-74%	4 of 4	0
	Gabriel Kreuther	21 of 40	37-67%	4 of 4	0
West Village	Buvette	36 of 40	77-96%	4 of 4	19
	Little Owl	26 of 40	50-78%	4 of 4	2
	Via Carota	19 of 40	33-63%	4 of 4	4
	L'Artusi	18 of 40	31-60%	4 of 4	7
Arthur Avenue	Zero Otto Nove	35 of 40	74-95%	4 of 4	18
	Mario’s Restaurant	28 of 40	55-82%	4 of 4	3
	Emilia's Restaurant	19 of 40	33-63%	4 of 4	2
	Full Moon Pizzeria	19 of 40	33-63%	4 of 4	0
St. George	Beso	35 of 40	74-95%	4 of 4	11
	Enoteca Maria	31 of 40	62-88%	4 of 4	9
	Ruddy & Dean	26 of 40	50-78%	4 of 4	3
	Pier 76	20 of 40	35-65%	3 of 4	0
Flushing	Asian Jewels *	23 of 40	42-71%	4 of 4	1
	Asian Jewels Seafood Restaurant	20 of 40	35-65%	4 of 4	6
	Hot Pot *	18 of 40	31-60%	4 of 4	0
	Jiang Nan	17 of 40	29-58%	3 of 4	4
Williamsburg	Roberta’s	23 of 40	42-71%	3 of 4	4
	L’Industrie Pizzeria	21 of 40	37-67%	4 of 4	14
	Best Pizza	17 of 40	29-58%	4 of 4	2
	Kellogg’s Diner	16 of 40	26-55%	4 of 4	3

Harlem	Red Rooster *	22 of 40	40-69%	4 of 4	0
	Red Rooster Harlem	18 of 40	31-60%	4 of 4	10
	Charles Pan-Fried Chicken	14 of 40	22-50%	4 of 4	2
	Clay	14 of 40	22-50%	4 of 4	3
Park Slope	Miriam	20 of 40	35-65%	4 of 4	12
	Stone Park Cafe	18 of 40	31-60%	4 of 4	5
	Palo Santo	15 of 40	24-53%	3 of 4	0
	Lore	14 of 40	22-50%	3 of 4	0
Upper West Side	Gray's Papaya	18 of 40	31-60%	4 of 4	16
	Essential by Christophe	12 of 40	18-45%	3 of 4	7
	Absolute Bagels	11 of 40	16-43%	3 of 4	2
	The Milling Room	11 of 40	16-43%	3 of 4	1

* 4 of the names above were matched in running text more often than they were written as a list entry. Treat them as doubtful until checked: some are neighbourhoods or cuisines, some are the short form of a longer name on the same list.

How the four assistants differ

Assistant	Answers	Cite a source	Names per answer	Distinct names	End mid-sentence
Claude	100	100 of 100	9.2	247	0 of 100
Gemini	100	35 of 100	8.0	205	0 of 100
ChatGPT	100	69 of 100	9.9	250	0 of 100
Perplexity	100	100 of 100	3.8	134	4 of 100

Names are counted only if they occur in at least 3 of the 400 answers. "End mid-sentence" is a check on the last character of the answer, a sign it may have been cut off.

Sites the assistants cite

Assistant	Most-cited sites (answers citing, of that assistant's answers)
Claude	yelp.com 57, theinfatuation.com 34, opentable.com 28, tripadvisor.com 28, timeout.com 20, exploringtheupperwestside.com 10
Gemini	reddit.com 25, facebook.com 14, nyctourism.com 10, timeout.com 9, resy.com 8, pier76si.com 7
ChatGPT	theinfatuation.com 31, opentable.com 29, reddit.com 29, tripadvisor.com 22, restaurantji.com 17, en.wikipedia.org 11
Perplexity	theinfatuation.com 86, opentable.com 82, tripadvisor.com 82, happy-hour.nyc 41, timeout.com 27, guide.michelin.com 23

A cited site is one the assistant returned with its answer. This report does not show that a citation is where a given name came from.

Read this before using any name

Name extraction is a heuristic and its error rate has not been measured. Three things are known about this list:

- 14 of the 96 names with 10 or more answers were matched in running text more often than they were written as a list entry. These are the first to distrust: Midtown East, Asian Jewels, Red Rooster, Hot Pot, Mountain House, Chela, Koreatown, Amylos, Korean BBQ, Greek / Mediterranean, Little Senegal, Prime Met Steakhouse, Sal, Kris & Charlie’s, Sage Thai.
- 48 pairs of counted names are a short and a long spelling of what may be one place in the same neighbourhood (for example AbuQir / AbuQir Seafood; Afuri Ramen / Afuri Ramen + Dumpling; Agnanti / Agnanti Meze). Each pair inflates the number of names; the short spelling's count already includes answers that used the long one.
- Checking sheet: 136 rows (all 96 names with 10 or more answers, plus 40 sampled from the 327 names with 3 to 9). Rows with a verdict: 0 of 136.

Decision rule

NOT EVALUATED. The rule set before the run needs a verdict on every row of the checking sheet (open, closed, not found, wrong area, or not a place). Until then there is no finding about whether a stable set of most-named restaurants exists in this market.

Predictions written down before the run

- NYC R3 (restaurants named by >=3 of 4 engines, same area): 40 (assistant-suggested, adopted by operator 2026-10-06)
- NYC R4 (top-10 share of verified mentions): 15 % (assistant-suggested, adopted by operator 2026-10-06)
- NYC R5 (>=10 tier, closed + not found + wrong area): 5 % (assistant-suggested, adopted by operator 2026-10-06)
- Engine with the most distinct restaurants named: ChatGPT (assistant-suggested, adopted by operator 2026-10-06)
- Market of the four with the smallest R3: Delaware (assistant-suggested, adopted by operator 2026-10-06)

They are defined on checked names, so they cannot be scored yet.

What this report does not say

- Nothing about how many diners any restaurant gets, or what share of its business comes from AI assistants. That needs each restaurant's own records.
- Nothing about which restaurant is better. A name given often is a name given often.
- Nothing about the consumer apps. These answers came from the developer API with a search tool offered; the apps add their own search, location and memory.
- Nothing beyond these 10 neighbourhoods, these twenty questions, and this one evening.

Check this yourself

Every answer is stored with the request that produced it and a SHA-256 hash of its text. The four evidence files sit in the repository beside this report. The checker is one Python file that needs nothing installed.

Evidence file	Clean rows	Hash match	Model as returned
multi_anthropic_nyc_restaurant_v1_20261006.jsonl	100 of 100	100 of 100	claude-sonnet-5
multi_gemini_nyc_restaurant_v1_20261006.jsonl	100 of 100	100 of 100	gemini-3.8-flash
multi_openai_nyc_restaurant_v1_20261006.jsonl	100 of 100	100 of 100	gpt-4o-2024-08-06
multi_perplexity_nyc_restaurant_v1_20261006.jsonl	100 of 100	100 of 100	perplexity/sonar

To check a file, run the line for it. The long value is the bundle digest of that file as this report computed it.

```
python verify_evidence.py multi_anthropic_nyc_restaurant_v1_20261006.jsonl --expect
a73dd2e498fca71d2302300290ee9f16cbfbe6465a544edb1c7b15591d43bf0b
python verify_evidence.py multi_gemini_nyc_restaurant_v1_20261006.jsonl --expect
81ff98610448ba70a6101e7f3ee3ef626503c8e92bfc272715c3ffff1d5d91ea
python verify_evidence.py multi_openai_nyc_restaurant_v1_20261006.jsonl --expect
674982a3ab7628faee1043c84792f22f1b49070b149b7fdeafb8d2f4e008c25b
python verify_evidence.py multi_perplexity_nyc_restaurant_v1_20261006.jsonl --expect
3eb6d15b25a23e025bf63c267dfae9821cb4653d36d5747a2487abc7a6294c85
```

A match shows the stored text and the stored hashes agree with each other and with this report. It does not show that the answers came from the assistants named, or when: that rests on the operator's run record.

Design fixed before the run

The questions, the counting rules, the decision rule and the analysis code were hashed and committed before the first answer was collected. Seal time as recorded in the seal file: 2026-10-07 02:25:44 UTC; commit 9deb12b33773. First answer collected: 2026-10-07 02:26:06 UTC.

Analysis cross-check: the 935 name counts in this report were recomputed from the evidence files and equal the committed analysis in results_nyc_v1.

1 character(s) outside the report font were reduced to their base letters: U+0113.