

AI VISIBILITY MEASUREMENT

Subject: **Credible PR** · generated 2026-09-01 22:05Z · make_pdf v0.1

This measures what AI assistants **answer**, not what people ask. It carries no information about search volume or purchase intent. Every figure is computed from the named evidence file at generation time; this document contains no typed numbers and no placeholders.

ENGINE	claude-sonnet-4-6
TEMPERATURE	0.0
PROMPTS	28 · 5 families
RUNS / PROMPT	5-7
CLEAN	191 of 210 attempted
RUNG	R1 - ran on the operator's machine

Credible PR appears in 0 of 191 responses.

95% Wilson interval on the overall rate: 0.0% - 2.0%. STATUS = COMPLETE for this evidence file. Per-family tables below print each rate against its own denominator; the bound travels with every zero so no claim can outrun the data.

BUYER_INTENT - n = 41 responses	Mention rate	95% CI
Reputation Ink	9.8% (4/41)	3.9% - 22.5%
Reputation Rhino	7.3% (3/41)	2.5% - 19.4%
Credible PR	0.0% (0/41)	0.0% - 8.6%
Clarity PR	0.0% (0/41)	0.0% - 8.6%
Brandstyle	0.0% (0/41)	0.0% - 8.6%

CATEGORY - n = 41 responses	Mention rate	95% CI
Clarity PR	53.7% (22/41)	38.7% - 67.9%
Reputation Rhino	19.5% (8/41)	10.2% - 34.0%
Reputation Ink	9.8% (4/41)	3.9% - 22.5%
Brandstyle	9.8% (4/41)	3.9% - 22.5%
Credible PR	0.0% (0/41)	0.0% - 8.6%

COMPARISON - n = 25 responses	Mention rate	95% CI
Reputation Ink	24.0% (6/25)	11.5% - 43.4%
Brandstyle	24.0% (6/25)	11.5% - 43.4%
Credible PR	0.0% (0/25)	0.0% - 13.3%
Clarity PR	0.0% (0/25)	0.0% - 13.3%
Reputation Rhino	0.0% (0/25)	0.0% - 13.3%

LONG_TAIL - n = 42 responses	Mention rate	95% CI
------------------------------	--------------	--------

Credible PR	0.0% (0/42)	0.0% - 8.4%
Clarity PR	0.0% (0/42)	0.0% - 8.4%
Reputation Ink	0.0% (0/42)	0.0% - 8.4%
Reputation Rhino	0.0% (0/42)	0.0% - 8.4%
Brandstyle	0.0% (0/42)	0.0% - 8.4%

PROBLEM - n = 42 responses	Mention rate	95% CI
Credible PR	0.0% (0/42)	0.0% - 8.4%
Clarity PR	0.0% (0/42)	0.0% - 8.4%
Reputation Ink	0.0% (0/42)	0.0% - 8.4%
Reputation Rhino	0.0% (0/42)	0.0% - 8.4%
Brandstyle	0.0% (0/42)	0.0% - 8.4%

Candidate list declared: Credible PR, Clarity PR, Reputation Ink, Reputation Rhino, Brandstyle. Nothing was scored off-list; entities outside this list are not measured by this document.

INSTRUMENT RECORD

Evidence file	pr_agency_run4.jsonl - one row per response: raw text, request payload, SHA-256
Prompt bank	pr_agency_v1 · 28 prompts · families: buyer_intent, category, comparison, long_tail, problem
Responses per family	41 · 41 · 25 · 42 · 42 - every rate prints against its own count
Rows attempted / clean / lost	210 / 191 / 19 (transport errors)
Runs per prompt	5, 6, 7 - uneven; per-family tables are the primary view
Surface declared	developer API · model claude-sonnet-4-6 · temperature 0.0. Not the consumer app surface; the gap is a stated limit, not pooled away.
Interval method	Wilson score, 95%, computed per rate at its own n
Bundle digest (SHA-256)	9eb756a2ca6ae2175bfe55a09d5b2c99b30e1451c2cd60050ddc0d1e2ee4abea

LIMITS - printed here, not discovered later

- One engine per model listed above (claude-sonnet-4-6). No claim is made about other engines or any consumer surface; adding them changes these numbers.
- Extraction error is unmeasured. Names are matched by a text rule with no published agreement coefficient yet. Every count depends on that unvalidated step; its measurement (an inter-labeller study on a sealed sample) is a prerequisite for any stronger language.
- One point in time. Models update; these rates decay. This file re-quoted next quarter without a re-run is a misquote.
- No link to revenue is claimed. Whether AI mention rates move sales is the open question this instrument is built to test - not one it has answered.

CHECK THIS YOURSELF

The evidence file ships with this report. The checker is public, standard library only, installs nothing:

```
git clone https://github.com/Rmasood1122/aivis-method.git
python3 aivis-method/verify_evidence.py pr_agency_run4.jsonl \
--expect 9eb756a2ca6ae2175bfe55a09d5b2c99b30e1451c2cd60050ddc0d1e2ee4abea
```

A PASS means the file was not altered after it was written - the record is tamper-evident and hash-verified. It does not mean collection was honest, or that these responses are what a consumer would see today. A hash cannot establish either, and this report says so rather than implying otherwise.